

ML-ZC Homework 2

```
In [8]: import pandas as pd
import ssl

# Disable SSL certificate verification
ssl._create_default_https_context = ssl._create_unverified_context

# Load the dataset from the URL
url = "https://raw.githubusercontent.com/alexeygrigorev/datasets/master/laptops.csv"
laptops_df = pd.read_csv(url)

# Display the first few rows to confirm
laptops_df.head()
```

```
Out[8]:
```

	Laptop	Status	Brand	Model	CPU	RAM	Storage	Storage type	GPU	Screen
0	ASUS ExpertBook B1 B1502CBA-EJ0436X Intel Core i5-1235U/16GB/512GB...	New	Asus	ExpertBook	Intel Core i5	8	512	SSD	NaN	15.6"
1	Alurin Go Start Intel Celeron N4020/8GB/256GB ...	New	Alurin	Go	Intel Celeron	8	256	SSD	NaN	15.6"
2	ASUS ExpertBook B1 B1502CBA-EJ0424X Intel Core i3-1215U/8GB/256GB...	New	Asus	ExpertBook	Intel Core i3	8	256	SSD	NaN	15.6"
3	MSI Katana GF66 12UC-082XES Intel Core i7-12700H/32GB/1000GB...	New	MSI	Katana	Intel Core i7	16	1000	SSD	RTX 3050	15.6"
4	HP 15S-FQ5085NS Intel Core i5-1235U/16GB/512GB...	New	HP	15S	Intel Core i5	16	512	SSD	NaN	15.6"

```
In [10]: # Update the column names by converting to lowercase and replacing spaces with underscores
laptops_df.columns = laptops_df.columns.str.lower().str.replace(' ', '_')

# Display the updated column names to confirm the changes
laptops_df.columns
```

```
Out[10]: Index(['laptop', 'status', 'brand', 'model', 'cpu', 'ram', 'storage',
               'storage_type', 'gpu', 'screen', 'touch', 'final_price'],
              dtype='object')
```

```
In [11]: # Select only the specified columns: 'ram', 'storage', 'screen', and 'final_price'
laptops_selected_columns = laptops_df[['ram', 'storage', 'screen', 'final_price']]

# Display the first few rows of the new dataframe to confirm
laptops_selected_columns.head()
```

Out[11]:

	ram	storage	screen	final_price
0	8	512	15.6	1009.00
1	8	256	15.6	299.00
2	8	256	15.6	789.00
3	16	1000	15.6	1199.00
4	16	512	15.6	669.01

```
In [13]: # First, standardize the column names (lowercase and replace spaces with underscores)
laptops_df.columns = laptops_df.columns.str.lower().str.replace(' ', '_')


# Now, try plotting the histogram again with the corrected column name
import matplotlib.pyplot as plt
import seaborn as sns

# Plot a histogram for the 'final_price' column
plt.figure(figsize=(8, 6))
sns.histplot(laptops_df['final_price'], bins=30, kde=True)
plt.title('Distribution of Final Price')
plt.xlabel('Final Price')
plt.ylabel('Frequency')
plt.show()


# Summary statistics (mean, median)
mean_price = laptops_df['final_price'].mean()
median_price = laptops_df['final_price'].median()

print(f"Mean: {mean_price}, Median: {median_price}")

# Box plot to detect outliers
plt.figure(figsize=(8, 6))
sns.boxplot(x=laptops_df['final_price'])
plt.title('Box Plot of Final Price')
plt.show()
```



Mean: 1312.638509259259, Median: 1031.9450000000002



Q1 Missing values

```
In [14]: # Check for missing values in each column
missing_values = laptops_df.isnull().sum()

# Display columns that have missing values
print(missing_values[missing_values > 0])
```

```
storage_type      42
gpu               1371
screen            4
dtype: int64
```

Q2 Median of RAM

```
In [16]: # Calculate the median of the 'ram' column (in lowercase)
ram_median = laptops_df['ram'].median()

print(f"The median (50th percentile) for RAM is: {ram_median} GB")
```

The median (50th percentile) for RAM is: 16.0 GB

```
In [17]: # Shuffle the dataset with a random seed of 42
shuffled_laptops_df = laptops_df.sample(frac=1, random_state=42).reset_index()

# Display the first few rows of the shuffled dataframe to confirm
shuffled_laptops_df.head()
```

```
Out[17]:
```

	laptop	status	brand	model	cpu	ram	storage	storage_type
0	MSI Modern 15 A11SBU-620XES Intel Core i7-1165...	Refurbished	MSI	Modern	Intel Core i7	32	1000	SSD
1	Thomson Neo White Intel Celeron N4020/4GB/64GB...	New	Thomson	Neo	Intel Celeron	4	64	eMMC
2	MSI Modern 14 B11SB-420XES Intel Core i7-1165G...	Refurbished	MSI	Modern	Intel Core i7	32	1000	SSD
3	ASUS ZenBook 13 OLED UX325EA Intel Evo Core i7...	New	Asus	ZenBook	Intel Evo Core i7	16	512	SSD
4	Acer Predator Triton 500 SE PT516-51s-75K6 Int...	New	Acer	Predator	Intel Core i7	32	1000	SSD

```
In [18]: from sklearn.model_selection import train_test_split

# First, split the data into training and temporary sets (60% train, 40% temp)
train_set, temp_set = train_test_split(laptops_df, test_size=0.4, random_state=42)

# Then, split the temporary set into validation and test sets (50% of 40% each)
val_set, test_set = train_test_split(temp_set, test_size=0.5, random_state=42)

# Display the sizes of each set
```

```
print(f"Training set size: {len(train_set)}")
print(f"Validation set size: {len(val_set)}")
print(f"Test set size: {len(test_set)}")
```

Training set size: 1296
Validation set size: 432
Test set size: 432

Question 3

```
In [21]: # Check the data types of the training dataset
print(train_set.dtypes)

# Select only numeric columns to calculate mean and fill missing values
numeric_columns = train_set.select_dtypes(include=['float64', 'int64']).columns

# Duplicate the training dataset
train_set_mean = train_set.copy() # For replacing with mean
train_set_zero = train_set.copy() # For replacing with zero

# Calculate the mean for each numeric column in the training set
mean_values = train_set_mean[numeric_columns].mean()

# Replace missing values in the first dataset with the mean values
train_set_mean[numeric_columns] = train_set_mean[numeric_columns].fillna(mean_values)

# Replace missing values in the second dataset with zero for numeric columns
train_set_zero[numeric_columns] = train_set_zero[numeric_columns].fillna(0)

# Optionally, display the first few rows to confirm changes
print("Train Set with Mean Values:")
print(train_set_mean.head())

print("\nTrain Set with Zero Values:")
print(train_set_zero.head())
```

Train Set with Mean Values:

		gpu	screen	touch	final_price
1357	RTX	3070	17.3	No	2228.59
1055		NaN	11.6	No	365.90
513	RTX	4050	14.0	No	1989.01
81	GTX	1650	15.6	No	999.00
1803	RTX	2060	15.6	No	1091.09

Train Set with Zero Values:

		gpu	screen	touch	final_price
1357	RTX	3070	17.3	No	2228.59
1055		NaN	11.6	No	365.90
513	RTX	4050	14.0	No	1989.01
81	GTX	1650	15.6	No	999.00
1803	RTX	2060	15.6	No	1091.09

file:///Users/sebastianreyes/Documents/3 Carrera/Proyectos/Portafolio/machinelearning_zoomcamp/ml_zc_h2.html

```

# Prepare the features (X) and target (y) for the models
# Assume 'final_price' is the target variable and the rest are features
features = ['ram', 'storage', 'screen'] # Adjust features as needed
target = 'final_price'

# Prepare X and y for the first model (mean values)
X_train_mean = train_set_mean[features]
y_train_mean = train_set_mean[target]
X_test = test_set[features] # Use the same features for the test set
y_test = test_set[target]

# Train the linear regression model with mean values
model_mean = LinearRegression()
model_mean.fit(X_train_mean, y_train_mean)

# Make predictions on the test set
y_pred_mean = model_mean.predict(X_test)

# Calculate RMSE for the model with mean replacement
rmse_mean = np.sqrt(mean_squared_error(y_test, y_pred_mean))

# Prepare X and y for the second model (zero values)
X_train_zero = train_set_zero[features]
y_train_zero = train_set_zero[target]

# Train the linear regression model with zero values
model_zero = LinearRegression()
model_zero.fit(X_train_zero, y_train_zero)

# Make predictions on the test set
y_pred_zero = model_zero.predict(X_test)

# Calculate RMSE for the model with zero replacement
rmse_zero = np.sqrt(mean_squared_error(y_test, y_pred_zero))

# Display the RMSE results
print(f"RMSE for the model with mean replacement: {rmse_mean}")
print(f"RMSE for the model with zero replacement: {rmse_zero}")

```

RMSE for the model with mean replacement: 684.07248803113

RMSE for the model with zero replacement: 684.089743558418

Question 4

```

In [23]: from sklearn.linear_model import Ridge
from sklearn.metrics import mean_squared_error
import numpy as np

# Prepare the features (X) and target (y) for the model with zero replacement
X_train_zero = train_set_zero[features]
y_train_zero = train_set_zero[target]

# Define the regularization strengths
r_values = [0, 0.01, 0.1, 1, 5, 10, 100]

# Store RMSE values for each r
rmse_results = []

# Train and evaluate Ridge regression for each r value
for r in r_values:

```

```

# Create and train the model
model_ridge = Ridge(alpha=r)
model_ridge.fit(X_train_zero, y_train_zero)

# Make predictions on the test set
y_pred_zero = model_ridge.predict(X_test)

# Calculate RMSE and round to two decimals
rmse_zero = np.sqrt(mean_squared_error(y_test, y_pred_zero))
rmse_results.append(round(rmse_zero, 2))

# Display the RMSE results for each r
for r, rmse in zip(r_values, rmse_results):
    print(f"RMSE for r={r}: {rmse}")

```

```

RMSE for r=0: 684.09
RMSE for r=0.01: 684.09
RMSE for r=0.1: 684.09
RMSE for r=1: 684.09
RMSE for r=5: 684.08
RMSE for r=10: 684.07
RMSE for r=100: 683.94

```

Question 5

```

In [25]: import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

# Initialize an empty list to collect RMSE scores
rmse_scores = []

# Define the seeds
seeds = [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]

# Prepare the features (X) and target (y) for the model
features = ['ram', 'storage', 'screen'] # Adjust this based on your dataset
target = 'final_price'

# Loop over each seed
for seed in seeds:
    # Split the dataset into train and test (60% train, 40% test)
    train_temp_set, test_set = train_test_split(laptops_df, test_size=0.4, r

    # Split the train_temp_set into train and validation (60% train, 40% test)
    train_set, val_set = train_test_split(train_temp_set, test_size=0.4, ran

    # Replace missing values with 0 in the training set
    train_set.fillna(0, inplace=True)

    # Replace missing values with 0 in the validation set to prevent NaN in
    val_set.fillna(0, inplace=True)

    # Prepare features and target for training and validation
    X_train = train_set[features]
    y_train = train_set[target]
    X_val = val_set[features]
    y_val = val_set[target]

```



```

# Train a linear regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Make predictions on the validation set
y_pred_val = model.predict(X_val)

# Calculate RMSE for the validation set
rmse_val = np.sqrt(mean_squared_error(y_val, y_pred_val))

# Append the RMSE score to the list
rmse_scores.append(rmse_val)

# Calculate the standard deviation of RMSE scores
std_rmse = np.std(rmse_scores)

# Round the standard deviation to 3 decimal places
std_rmse_rounded = round(std_rmse, 3)

# Display the standard deviation of RMSE scores
print(f"Standard deviation of RMSE scores: {std_rmse_rounded}")

```

Standard deviation of RMSE scores: 25.473

Question 6

```

In [26]: import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Ridge
from sklearn.metrics import mean_squared_error

# Prepare the features (X) and target (y) for the model
features = ['ram', 'storage', 'screen'] # Adjust this based on your dataset
target = 'final_price'

# Split the dataset into train and test (60% train, 40% test) using seed 9
train_temp_set, test_set = train_test_split(laptops_df, test_size=0.4, random_state=9)

# Split the train_temp_set into train and validation (60% train, 40% temp = 24% total)
train_set, val_set = train_test_split(train_temp_set, test_size=0.4, random_state=9)

# Combine train and validation datasets
combined_train_set = pd.concat([train_set, val_set])

# Replace missing values with 0 in the combined dataset
combined_train_set.fillna(0, inplace=True)

# Prepare features and target for training
X_train = combined_train_set[features]
y_train = combined_train_set[target]

# Prepare features for the test set (replace missing values with 0)
test_set.fillna(0, inplace=True)
X_test = test_set[features]
y_test = test_set[target]

# Train a Ridge regression model with r = 0.001
model = Ridge(alpha=0.001)
model.fit(X_train, y_train)

```

```
# Make predictions on the test set
y_pred_test = model.predict(X_test)

# Calculate RMSE for the test set
rmse_test = np.sqrt(mean_squared_error(y_test, y_pred_test))

# Display the RMSE for the test dataset
print(f"RMSE on the test dataset: {rmse_test:.3f}")
```

RMSE on the test dataset: 612.863